

# Lab 1

## *IMC 490: Machine Learning for IMC*

In this lab, we'll be going over the following topics:

- operations and data structures
- dataframes
- accessing data
- indexing data
- data types
- gotchas
- regression

Website for Introduction to Statistical Learning: <http://www-bcf.usc.edu/~gareth/ISL/>

**Remember to use `?command` or `help(command)` in the R console to access documentation at any time if you have questions.**

### Navigation

```
setwd()
getwd()
list.files()
```

The above three commands will allow you to navigate your filesystem. You'll use `getwd()` (**get working directory**) to figure out where you are, `setwd()` (**set working directory**) to move around, and `list.files()` to look at the stuff that's in the folder.

```
setwd("~/Documents/Machine-Learning-IMC490/Lab1/")
```

```
getwd()
```

```
## [1] "/home/eric/Documents/Machine-Learning-IMC490/Lab1"
```

```
list.files()
```

```
## [1] "HW1"      "Lab1.pdf" "Lab1.Rmd"
```

**Common gotcha:** Be sure to enter the filepath in `setwd()` as a string (in quotes " "). A common mistake is to forget the quotes. If you do, R will think that `/your/filepath/` is a variable holding some value... And complain when it finds that it isn't.

## Data Types and Vectors

`c()`

From the R documentation: “R has six basic (‘atomic’) vector types: logical, integer, real, complex, string (or character) and raw.”

We often use integer, real, string, and logical types. In normal use we almost never see complex and raw.

Let’s make some vectors. `c()` (combine) is the generic function for creating a vector.

```
# 2 ways of making a vector with numbers 1 through 5  
c(1, 2, 3, 4, 5)
```

```
## [1] 1 2 3 4 5
```

```
c(1:5)
```

```
## [1] 1 2 3 4 5
```

```
# create a character (string) vector and check its type  
a_char_vec = c("a", "b", "c")  
typeof(a_char_vec)
```

```
## [1] "character"
```

```
# create a boolean vector  
some_numbers = c(1:5)  
a_bool_vec = some_numbers >= 3  
a_bool_vec
```

```
## [1] FALSE FALSE TRUE TRUE TRUE
```

The operation `>= 3` was vectorized and applied to each element of `some_numbers`. Vectorized operations are at the core of R. A vectorized operation is an operation that is applied to each element of a vector. This includes arithmetic and comparison operations.

```
# create a vector 1:10 and add 5 to each element  
x = c(1:10)  
x + 5
```

```
## [1] 6 7 8 9 10 11 12 13 14 15
```

**Common gotcha:** If you perform an operation with two vectors of different length, the shorter vector will be extended to complete the operation.

```
x = c(1, 2, 3, 4, 5, 6, 7, 8, 9, 10)  
y = c(0, 100)  
x + y
```

```
## [1] 1 102 3 104 5 106 7 108 9 110
```

**Some more vector functions:**

Casting vectors from one type to another:

```
as.integer()  
as.numeric()  
as.character()
```

Checking characteristics of a vector: `typeof()`

```
length()  
unique()
```

Calculating statistics of a numeric vector:

```
mean()  
sd()  
min()  
max()
```

## Dataframes

```
data.frame()
```

A dataframe is simply a collection of vectors. Here we're making a dataframe with column x which is an index, and y, which are random samples from a standard normal distribution.

```
dat = data.frame(x = c(1:10), y = rnorm(10))
dat
```

```
##      x      y
## 1    1 -0.54881116
## 2    2  0.73988377
## 3    3 -0.11700799
## 4    4  0.08545879
## 5    5 -1.44733763
## 6    6  0.33142182
## 7    7 -0.54676942
## 8    8  1.09998854
## 9    9  0.75496501
## 10 10 -0.31216168
```

To extract vectors from a dataframe, use the dollar sign operator.

```
dat$y
```

```
## [1] -0.54881116  0.73988377 -0.11700799  0.08545879 -1.44733763
## [6]  0.33142182 -0.54676942  1.09998854  0.75496501 -0.31216168
```

## Reading in data

```
read.csv() data()
```

Link to sample dataset: <http://www-bcf.usc.edu/~gareth/ISL/Auto.csv>